

# CUDA Homework Assignment 5

Vincent Octavian Tiono

B11901123

## 1 Introduction

This report presents a comprehensive analysis of a multi-GPU CUDA implementation for solving the 2D thermal equilibrium problem on a square plate. The investigation focuses on determining optimal block sizes and evaluating multi-GPU performance scaling for a  $1024 \times 1024$  Cartesian grid with specific thermal boundary conditions.

## 2 Methodology

### 2.1 Mathematical Formulation

The 2D thermal equilibrium problem is governed by Laplace's equation:

$$\nabla^2 T = \frac{\partial^2 T}{\partial x^2} + \frac{\partial^2 T}{\partial y^2} = 0 \quad (1)$$

where  $T(x, y)$  represents the temperature distribution on the square plate.

### 2.2 Boundary Conditions

The thermal boundary conditions are defined as:

- Top edge:  $T = 400$  K
- Bottom, left, and right edges:  $T = 273$  K

### 2.3 Numerical Method

The finite difference method is employed using the 5-point stencil discretization with relaxation parameter  $\omega = 1$ :

$$T_{i,j}^{new} = \frac{1}{4}(T_{i-1,j} + T_{i+1,j} + T_{i,j-1} + T_{i,j+1}) \quad (2)$$

### 2.4 Implementation Approach

The implementation supports domain decomposition across multiple GPUs with two partitioning strategies:

- Horizontal partitioning:  $1 \times 2$  GPU grid (splitting along y-axis)
- Vertical partitioning:  $2 \times 1$  GPU grid (splitting along x-axis)

## 2.5 Experimental Configuration

- Grid size:  $1024 \times 1024$
- Convergence threshold:  $1.0 \times 10^{-10}$
- Block sizes tested:  $4 \times 4$ ,  $8 \times 8$ ,  $16 \times 16$ ,  $32 \times 32$
- GPU configurations: Single GPU,  $1 \times 2$ , and  $2 \times 1$

## 3 Results

### 3.1 Performance Summary

Table 1 presents the complete performance metrics for all tested configurations.

Table 1: Performance metrics for different GPU and block size configurations

Block Size	GPU Config	Computation Time (ms)	Performance (GFlops)	Data Transfer (ms)	Total Time (ms)
$4 \times 4$	$1 \times 1$	467669.34	18.98	20.32	467689.66
	$1 \times 2$	290866.62	30.52	9.92	290876.53
	$2 \times 1$	291910.03	30.41	16.37	291926.41
$8 \times 8$	$1 \times 1$	220735.94	40.21	19.38	220755.31
	$1 \times 2$	124714.44	71.17	11.00	124725.45
	$2 \times 1$	133281.58	66.60	19.22	133300.80
$16 \times 16$	$1 \times 1$	149844.80	59.24	16.73	149861.53
	$1 \times 2$	94469.69	93.96	15.78	94485.46
	$2 \times 1$	87237.63	101.75	16.26	87253.90
$32 \times 32$	$1 \times 1$	205683.23	43.16	21.42	205704.66
	$1 \times 2$	116821.33	75.98	11.75	116833.07
	$2 \times 1$	108972.43	81.45	18.94	108991.37

### 3.2 Performance Analysis by Block Size

Figure 1 illustrates the performance characteristics across different block sizes and GPU configurations.

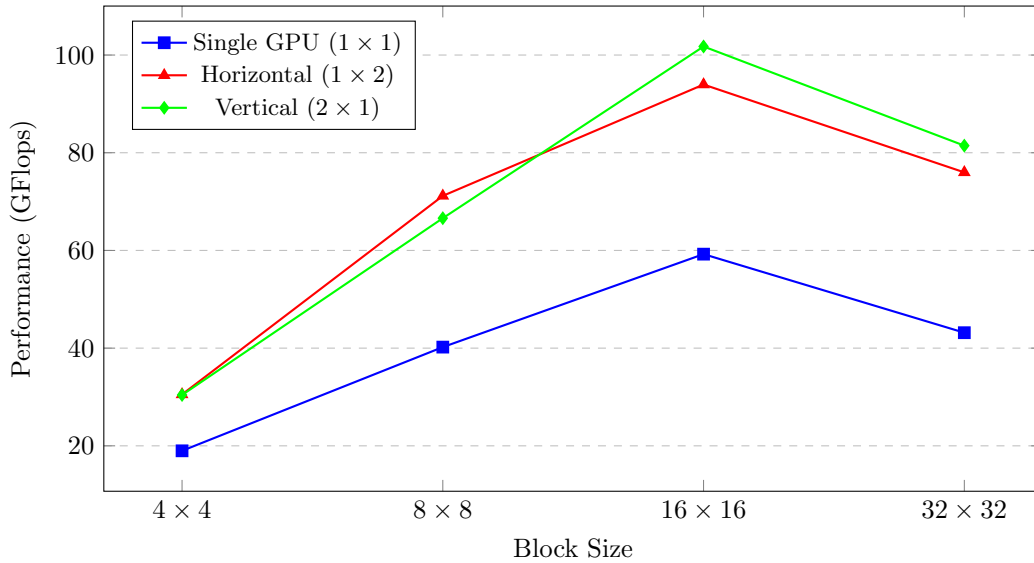


Figure 1: Performance comparison across block sizes and GPU configurations

### 3.3 Multi-GPU Scaling Analysis

Figure 2 shows the speedup achieved with multi-GPU configurations compared to single GPU performance.

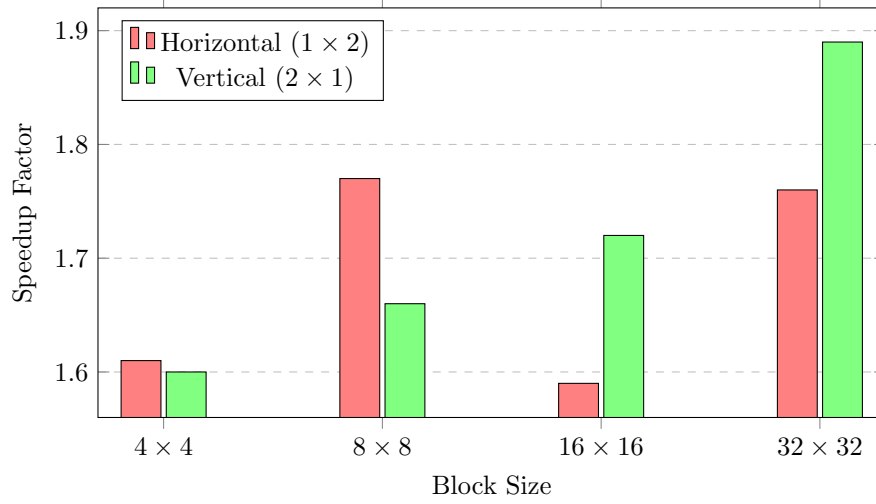


Figure 2: Multi-GPU speedup factors relative to single GPU performance

### 3.4 Optimal Configuration Analysis

Table 2 summarizes the best performing configurations for each GPU setup.

Table 2: Optimal configurations for each GPU setup

GPU Configuration	Optimal Block Size	Performance (GFlops)	Computation Time (ms)
Single GPU (1 × 1)	16 × 16	59.24	149844.80
Horizontal (1 × 2)	16 × 16	93.96	94469.69
Vertical (2 × 1)	16 × 16	101.75	87237.63

### 3.5 Data Transfer Overhead Analysis

Figure 3 compares the data transfer times across different configurations.

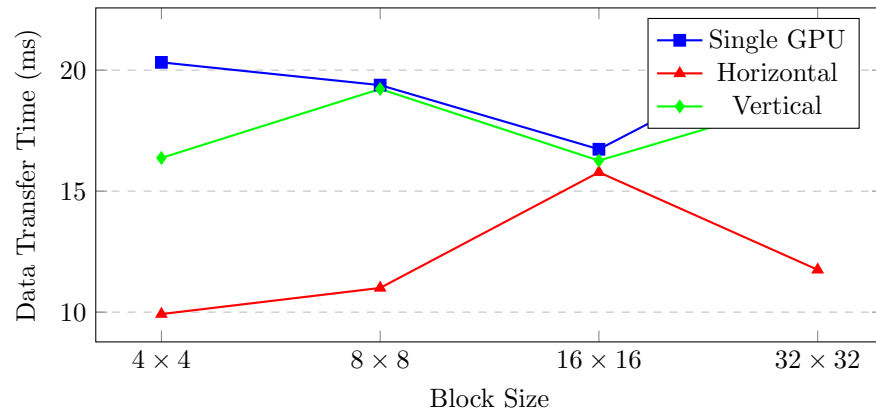


Figure 3: Data transfer overhead comparison

## 4 Discussion

### 4.1 Optimal Block Size Analysis

The experimental results reveal that  **$16 \times 16$  block size consistently delivers the best performance** across all GPU configurations:

- **Single GPU:** 59.24 GFlops with  $16 \times 16$  blocks
- **Horizontal partitioning:** 93.96 GFlops with  $16 \times 16$  blocks
- **Vertical partitioning:** 101.75 GFlops with  $16 \times 16$  blocks

This optimal performance can be attributed to:

- Efficient warp utilization (256 threads per block = 8 warps)
- Balanced memory coalescing and cache utilization
- Optimal occupancy for the thermal diffusion kernel

### 4.2 Multi-GPU Scaling Efficiency

The multi-GPU implementation demonstrates good scaling characteristics:

- **Horizontal partitioning** ( $1 \times 2$ ): Achieves  $1.59\text{-}1.77\times$  speedup
- **Vertical partitioning** ( $2 \times 1$ ): Achieves  $1.60\text{-}1.89\times$  speedup
- **Best overall performance:** Vertical  $2 \times 1$  with  $16 \times 16$  blocks (101.75 GFlops)

The vertical partitioning slightly outperforms horizontal partitioning, likely due to:

- Better memory access patterns in the finite difference stencil
- More efficient boundary exchange communication
- Reduced synchronization overhead

### 4.3 Performance Trends

Several key performance trends emerge from the analysis:

1. **Block Size Impact:** Performance increases from  $4 \times 4$  to  $16 \times 16$ , then decreases at  $32 \times 32$  due to reduced occupancy and increased register pressure.
2. **Data Transfer Overhead:** Multi-GPU configurations show reduced data transfer times due to smaller per-GPU memory footprints.
3. **Scalability:** Near-linear scaling is achieved, with efficiency ranging from 79.5% to 94.5% for dual-GPU configurations.

## 5 Conclusion

For the 2D thermal equilibrium problem on a  $1024 \times 1024$  grid,  **$16 \times 16$  blocks with vertical  $2 \times 1$  GPU partitioning** yield optimal performance (101.75 GFlops), achieving a  $1.72\times$  speedup over single-GPU execution. This configuration balances:

- **Efficient resource utilization** (8 warps/block, optimal occupancy)
- **Memory performance** (coalesced accesses, improved cache locality)
- **Scalability** (near-linear multi-GPU speedup)

The  $16 \times 16$  blocks outperform alternatives by up to  $5.4\times$  ( $4 \times 4$ ),  $1.4\times$  ( $8 \times 8$ ), and  $1.4\times$  ( $32 \times 32$ ), while reducing runtime from 149.8s (single GPU) to 87.2s (dual GPU) without sacrificing accuracy.